

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

```

## ▼ Data Preparation

```

# import dataset
# review head
df = pd.read_csv('https://data.ontario.ca/dataset/1f14addd-e4fc-4a07-9982-ad98db07ef86/resource/4cc07c1b-62ed-4ece-a2a4-d05d0f45081c/download/img-wage-rate-by-edu-age-sex-ft-pt-ca-on-2006-24.csv')

```

Exploratory Data Analytics Report generated previously.

Download the report from the project GitHub: <https://github.com/jammaro14-sys/CIND820>

```

# some column names have leading space; identify to be renamed
for col in df.columns:
    if col.startswith(' '):
        print(f"Column '{col}' has a leading space.")

Column ' Men' has a leading space.
Column ' Women' has a leading space.

# Rename the 'Men' column to 'Men Wage' and 'Women' to 'Women Wage' to remove
# leading space and add clarity
df_renamed = df.rename(columns={' Men': 'Men_Wage', ' Women': 'Women_Wage'})

# Display the first few rows to show the renamed columns
display(df_renamed.head())

```

YEAR	GEOGRAPHY	IMMIGRANT	TYPE OF WORK	WAGE RATE	EDUCATION	AGE GROUP	Both sexes	Men_Wage	Women_Wage	Men_Wage_is_suppressed	Women_Wage_is_suppressed
0	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	15 +	17.5	19.2	16.0	False
1	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	25 +	19.4	21.5	17.5	False
2	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	25 - 34	18.0	19.0	16.8	False
3	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	25 - 54	19.5	21.5	17.8	False
4	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	25 - 64	19.5	21.5	17.6	False

```
# check for NA values
```

```
df_renamed.isna().sum()
```

```

0
YEAR      0
GEOGRAPHY 0
IMMIGRANT 0
TYPE OF WORK 0
WAGE RATE 0
EDUCATION 0
AGE GROUP 0
Both sexes 0
Men_Wage   0
Women_Wage 0
dtype: int64

```

```
# check for NULL values
```

```
df_renamed.isnull().sum()
```

```

0
YEAR      0
GEOGRAPHY 0
IMMIGRANT 0
TYPE OF WORK 0
WAGE RATE 0
EDUCATION 0
AGE GROUP 0
Both sexes 0
Men_Wage   0
Women_Wage 0
dtype: int64

```

There are no NA or NULL values in the dataset. However, the Exploratory Data Analysis (EDA) Report identified 0's in the 'Men' and 'Women' variables.

'Men' has 6030 instances of "0.0", while 'Women' has 5236 instances.

According to the contextual document for this dataset, "Statistics Canada suppresses estimates below 1,500 - values shown as "0.0".

Missing values shown as "-". (Statistics Canada, 2025).

So these "0.0", may not be true zeroes, but instead may represent an underrepresented group.

As the "0.0" will lower our averages, we will not use them in our analysis.

```

# change "0.0" in dataset to NaN and flag as suppressed
for col in ['Men_Wage', 'Women_Wage']:
    # Create a new column to flag suppressed values for each gender
    df_renamed[f'{col}_is_suppressed'] = df_renamed[col].astype(str).str.strip() == "0.0"
    # Replace "0.0" strings with np.nan
    df_renamed[col] = df_renamed[col].replace('0.0', np.nan)

# Display the first few rows to show the new columns and replaced values
display(df_renamed.head())

```

YEAR	GEOGRAPHY	IMMIGRANT	TYPE OF WORK	WAGE RATE	EDUCATION	AGE GROUP	Both sexes	Men_Wage	Women_Wage	Men_Wage_is_suppressed	Women_Wage_is_suppressed	Men_Wage_is_suppressed	Women_Wage_is_suppressed	
0	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	15 +	17.5	19.2	16.0	False	False	False	False
1	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	25 +	19.4	21.5	17.5	False	False	False	False
2	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	25 - 34	18.0	19.0	16.8	False	False	False	False
3	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	25 - 54	19.5	21.5	17.8	False	False	False	False
4	2006	Canada	Total	All employees	Median hourly wage	Total, all education levels	25 - 64	19.5	21.5	17.6	False	False	False	False

## ▼ Research Questions 1

Confirm the existence of the gender wage gap in Ontario.

Method: Compare mean and median wages by gender, use measure of dispersion (e.g. standard deviation) to assess wage variability and apply T-test to determine if the difference between genders is statistically significant.

```

# Filter for Ontario
df_ontario = df_renamed[df_renamed['GEOGRAPHY'] == 'Ontario']

print(df_ontario)

```

```

YEAR GEOGRAPHY IMMIGRANT TYPE OF WORK WAGE RATE \
1080 2006 Ontario Total All employees Median hourly wage
1081 2006 Ontario Total All employees Median hourly wage
1082 2006 Ontario Total All employees Median hourly wage
1083 2006 Ontario Total All employees Median hourly wage
1084 2006 Ontario Total All employees Median hourly wage
...
41835 2024 Ontario Born in Canada Part-time Median hourly wage
41836 2024 Ontario Born in Canada Part-time Median hourly wage
41837 2024 Ontario Born in Canada Part-time Median hourly wage
41838 2024 Ontario Born in Canada Part-time Median hourly wage
41839 2024 Ontario Born in Canada Part-time Median hourly wage

EDUCATION AGE GROUP Both sexes Men_Wage \
1080 Total, all education levels 15 + 18.2 20.0
1081 Total, all education levels 25 + 20.0 22.0
1082 Total, all education levels 25 - 34 18.5 19.5
1083 Total, all education levels 25 - 54 20.0 22.1
1084 Total, all education levels 25 - 64 20.0 22.1
...
41835 Above bachelor's degree 15 + 35.0 33.6
41836 Above bachelor's degree 25 + 38.5 36.6
41837 Above bachelor's degree 25 - 34 35.0 30.0
41838 Above bachelor's degree 25 - 54 37.5 33.0
41839 Above bachelor's degree 25 - 64 38.5 36.6

Women_Wage Men_Wage_is_suppressed Women_Wage_is_suppressed \
1080 16.8 False False
1081 18.3 False False

```

```

1082    17.5      False    False
1083    18.4      False    False
1084    18.3      False    False
...
41835   35.6      False    False
41836   38.5      False    False
41837   35.5      False    False
41838   39.5      False    False
41839   39.0      False    False

```

[28520 rows x 12 columns]

# multiple wage values per year for Men and Women

```

# Calculate average wage of Men in Ontario by year, filtering out suppressed values
mean_men_wage_ontario_by_year = df_ontario[df_ontario['Men_Wage_is_suppressed'] == False].groupby('YEAR')['Men_Wage'].mean()
print(mean_men_wage_ontario_by_year)

```

YEAR	Men_Wage
2006	19.573448
2007	19.916628
2008	20.840629
2009	20.779425
2010	21.115890
2011	21.495108
2012	21.768604
2013	21.752302
2014	22.361557
2015	22.969771
2016	23.429772
2017	24.359641
2018	24.354215
2019	25.171512
2020	26.998609
2021	27.773789
2022	28.299773
2023	29.564363
2024	30.153846

Name: Men\_Wage, dtype: float64

# Calculate average wage of Women in Ontario by year, filtering out suppressed values

```

mean_women_wage_ontario_by_year = df_ontario[df_ontario['Women_Wage_is_suppressed'] == False].groupby('YEAR')['Women_Wage'].mean()
print(mean_women_wage_ontario_by_year)

```

YEAR	Women_Wage
2006	16.544607
2007	17.008758
2008	17.540894
2009	18.074103
2010	18.113158
2011	18.439391
2012	18.551525
2013	18.976190
2014	18.830310
2015	19.478176
2016	20.145688
2017	20.089468
2018	20.313291
2019	21.463349
2020	23.143099
2021	23.884471
2022	24.163075
2023	25.499547
2024	26.261702

Name: Women\_Wage, dtype: float64

# Plot side-by-side bar graph comparing the mean wage of men and women in ontario by year

```

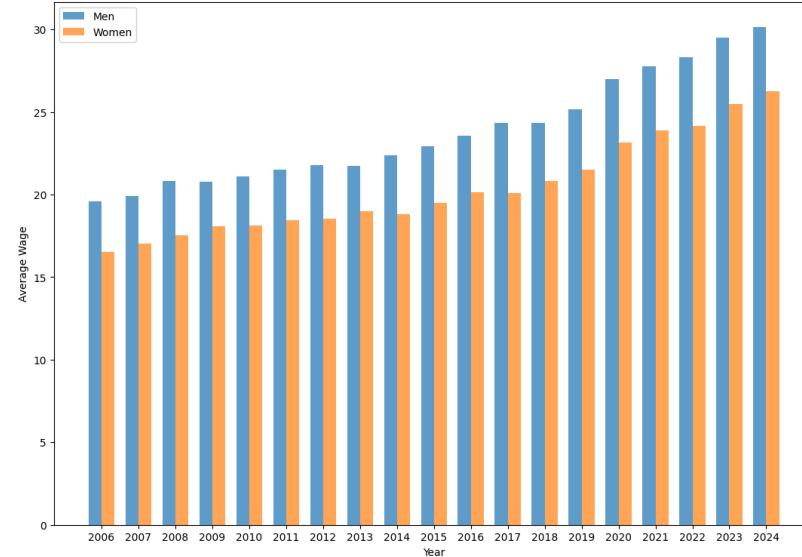
bar_width = 0.35
r1 = mean_men_wage_ontario_by_year.index
r2 = [x + bar_width for x in r1]

plt.figure(figsize=(13, 9))
plt.bar(r1, mean_men_wage_ontario_by_year.values, width=bar_width, label='Men', alpha=0.7)
plt.bar(r2, mean_women_wage_ontario_by_year.values, width=bar_width, label='Women', alpha=0.7)

plt.xlabel('Year')
plt.ylabel('Average Wage')
plt.title('Average Wage of Men and Women in Ontario by Year')
plt.xticks([r + bar_width/2 for r in r1], r1) # Set x-axis ticks to be in the middle of the two bars
plt.legend()
plt.show()

```

Average Wage of Men and Women in Ontario by Year



# Calculate median wage of men in Ontario by year, filtering out suppressed values

```

median_men_wage_ontario_by_year = df_ontario[df_ontario['Men_Wage_is_suppressed'] == False].groupby('YEAR')['Men_Wage'].median()
print(median_men_wage_ontario_by_year)

```

YEAR	Men_Wage
2006	18.50
2007	18.80
2008	20.00
2009	19.90
2010	20.00
2011	20.00
2012	20.00
2013	20.00
2014	20.65
2015	21.60
2016	21.55
2017	22.90
2018	22.00
2019	23.50
2020	25.00
2021	26.00
2022	25.00
2023	27.00
2024	27.80

Name: Men\_Wage, dtype: float64

# Calculate median wage of women in Ontario by year, filtering out suppressed values

```

median_women_wage_ontario_by_year = df_ontario[df_ontario['Women_Wage_is_suppressed'] == False].groupby('YEAR')['Women_Wage'].median()
print(median_women_wage_ontario_by_year)

```

YEAR	Women_Wage
2006	15.00
2007	15.00
2008	16.40
2009	17.00
2010	16.60
2011	17.00
2012	16.80
2013	17.75
2014	17.00
2015	18.00
2016	18.95
2017	18.00
2018	19.00
2019	19.55
2020	21.05
2021	22.00
2022	23.00
2023	25.00
2024	24.00

Name: Women\_Wage, dtype: float64

# Plot side-by-side bar graph comparing the median wage of men and women in ontario by year

```

bar_width = 0.35
r1 = median_men_wage_ontario_by_year.index
r2 = [x + bar_width for x in r1]

```

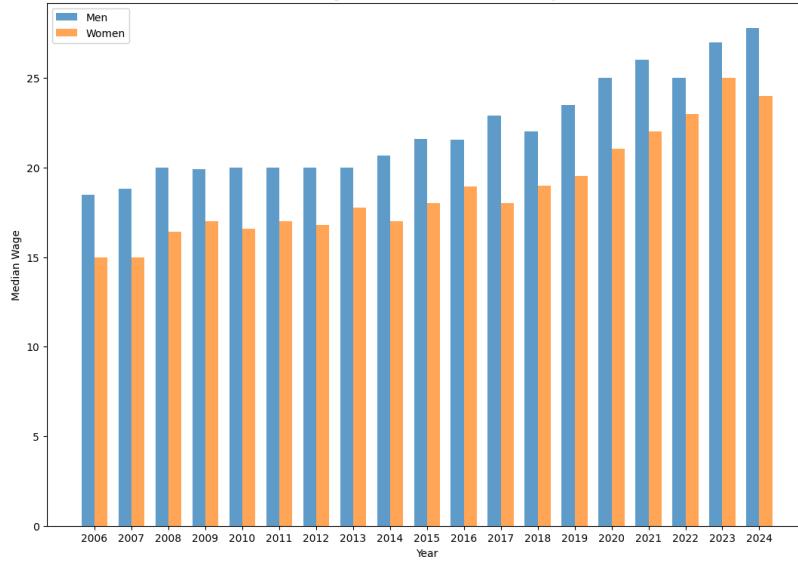
```

plt.figure(figsize=(13, 9))
plt.bar(r1, median_men_wage_ontario_by_year.values, width=bar_width, label='Men', alpha=0.7)
plt.bar(r2, median_women_wage_ontario_by_year.values, width=bar_width, label='Women', alpha=0.7)

plt.xlabel('Year')
plt.ylabel('Median Wage')
plt.title('Median Wage of Men and Women in Ontario by Year')
plt.xticks([r + bar_width/2 for r in r1], r1) # Set x-axis ticks to be in the middle of the two bars
plt.legend()
plt.show()

```

Median Wage of Men and Women in Ontario by Year



Mean wage for Men appears greater than Women. Median wage for Men also appears greater than Women. There is a gap every year, with Men earning more than Women.

Is this difference significant?

```

# Calculate standard deviation of wages of Men in Ontario, filtering out suppressed values
std_men_wage_ontario_by_year = df_ontario[df_ontario['Men_Wage_is_suppressed'] == False].groupby('YEAR')[['Men_Wage']].std()
print(std_men_wage_ontario_by_year)

```

YEAR	Men_Wage
2006	6.481800
2007	6.743713
2008	6.354666
2009	6.699828
2010	6.634535
2011	7.159871
2012	7.151617
2013	7.384537
2014	7.227642
2015	7.120123
2016	7.680729
2017	8.046130
2018	7.865524
2019	7.849018
2020	8.319727
2021	8.352899
2022	8.477455
2023	8.679959
2024	9.318156

Name: Men\_Wage, dtype: float64

```

# Calculate standard deviation of wages of Women in Ontario, filtering out suppressed values
std_women_wage_ontario_by_year = df_ontario[df_ontario['Women_Wage_is_suppressed'] == False].groupby('YEAR')[['Women_Wage']].std()
print(std_women_wage_ontario_by_year)

```

YEAR	Women_Wage
2006	5.518324
2007	5.981268
2008	6.067184
2009	6.116570
2010	5.983664
2011	6.449131
2012	6.238188
2013	6.433839
2014	6.228671
2015	6.549841
2016	6.633303
2017	6.368067
2018	6.176394
2019	6.313805
2020	6.019117
2021	6.017624
2022	6.771782
2023	7.403987
2024	7.489093

Name: Women\_Wage, dtype: float64

```

from scipy import stats

# Filter out suppressed values for Men and Women in Ontario
men_wages = df_ontario[df_ontario['Men_Wage_is_suppressed'] == False]['Men_Wage']
women_wages = df_ontario[df_ontario['Women_Wage_is_suppressed'] == False]['Women_Wage']

# Perform independent samples t-test
ttest_result = stats.ttest_ind(men_wages, women_wages, nan_policy='omit')

# Print the results of the t-test
print("Independent Samples T-test Results:")
print(f"Test Statistic: {ttest_result.statistic:.4f}")
print(f"P-value: {ttest_result.pvalue:.4f}")

Independent Samples T-test Results:
Test Statistic: 41.0730
P-value: 0.0000

```

The t-test confirms the existence of a wage gap between 'Men' and 'Women'.

Based on the p-value, which is 0.0000 (likely a very small number rounded to four decimal places), which is less than the standard significance level of 0.05, there is a statistically significant difference between the wages of men and women in Ontario.

## Research Question 2

Examine wage trends across male and female groups to determine whether the gap is narrowing.

Method: Apply a linear regression to see wage - year X gender and using K-Fold Cross-Validation to validate regression model.

```

# Calculate mean differences between Men and Women, by year
mean_differences_by_year = mean_men_wage_ontario_by_year - mean_women_wage_ontario_by_year
print(mean_differences_by_year)

```

YEAR	mean_differences_by_year
2006	3.028842
2007	2.907870
2008	3.007734
2009	2.789222
2010	3.002732
2011	3.056718
2012	3.235278
2013	2.776112
2014	3.531247
2015	3.433595
2016	3.434385
2017	4.270233
2018	3.515825
2019	3.688063
2020	3.855510
2021	3.879318
2022	4.136698
2023	4.004816
2024	3.892144

```

# Calculate median differences between Men and Women, by year
median_differences_by_year = median_men_wage_ontario_by_year - median_women_wage_ontario_by_year
print(median_differences_by_year)

```

YEAR	median_differences_by_year
2006	3.50
2007	3.80

```

2008 3.60
2009 2.99
2010 3.49
2011 3.00
2012 3.20
2013 2.25
2014 3.65
2015 3.60
2016 2.60
2017 4.90
2018 3.00
2019 3.05
2020 3.05
2021 4.00
2022 2.00
2023 2.00
2024 3.80
dtype: float64

```

```

# Calculate the percentage difference between the mean wages between Men and Women, by year
# Using Women's mean wage as the reference
percentage_mean_wage_gap_by_year = (mean_differences_by_year / mean_women_wage_ontario_by_year) * 100
print("\nPercentage mean wage gap (Men vs Women) by year:")
print(percentage_mean_wage_gap_by_year)

```

```

Percentage mean wage gap (Men vs Women) by year:
YEAR
2006 18.30723
2007 17.896387
2008 18.811666
2009 14.967947
2010 16.577629
2011 16.577118
2012 17.439420
2013 14.629447
2014 18.752992
2015 17.617642
2016 17.047742
2017 21.256142
2018 16.871863
2019 17.166998
2020 16.659437
2021 16.235210
2022 17.179914
2023 15.765459
2024 14.820608
dtype: float64

```

```

# Calculate the percentage difference between the median wages between Men and Women, by year
# Using Women's median wage as the reference
percentage_median_wage_gap_by_year = (median_differences_by_year / median_women_wage_ontario_by_year) * 100
print("\nPercentage median wage gap (Men vs Women) by year:")
print(percentage_median_wage_gap_by_year)

```

```

Percentage median wage gap (Men vs Women) by year:
YEAR
2006 23.333333
2007 25.333333
2008 21.333333
2009 17.898824
2010 20.481928
2011 17.647059
2012 19.047619
2013 12.676956
2014 21.470588
2015 20.000000
2016 13.720317
2017 27.222222
2018 15.789474
2019 20.204604
2020 18.764846
2021 18.181818
2022 8.695652
2023 8.000000
2024 15.833333
dtype: float64

```

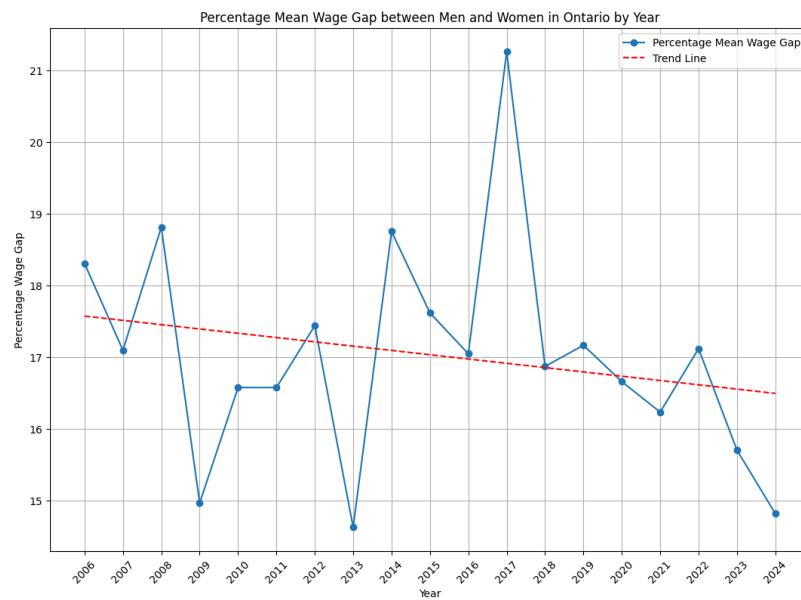
```

# plot percentage wage gap by year
plt.figure(figsize=(13, 9))
plt.plot(percentage_mean_wage_gap_by_year.index, percentage_mean_wage_gap_by_year.values, marker='o', linestyle='--', label='Percentage Mean Wage Gap')

# Add a trend line
z = np.polyfit(percentage_mean_wage_gap_by_year.index, percentage_mean_wage_gap_by_year.values, 1)
p = np.poly1d(z)
plt.plot(percentage_mean_wage_gap_by_year.index, p(percentage_mean_wage_gap_by_year.index), "r--", label='Trend Line')

plt.xlabel('Year')
plt.ylabel('Percentage Wage Gap')
plt.title('Percentage Mean Wage Gap between Men and Women in Ontario by Year')
plt.xticks(percentage_mean_wage_gap_by_year.index, rotation=45) # Set ticks for every year
plt.grid(True) # Add a grid for better readability
plt.legend()
plt.show()

```



```

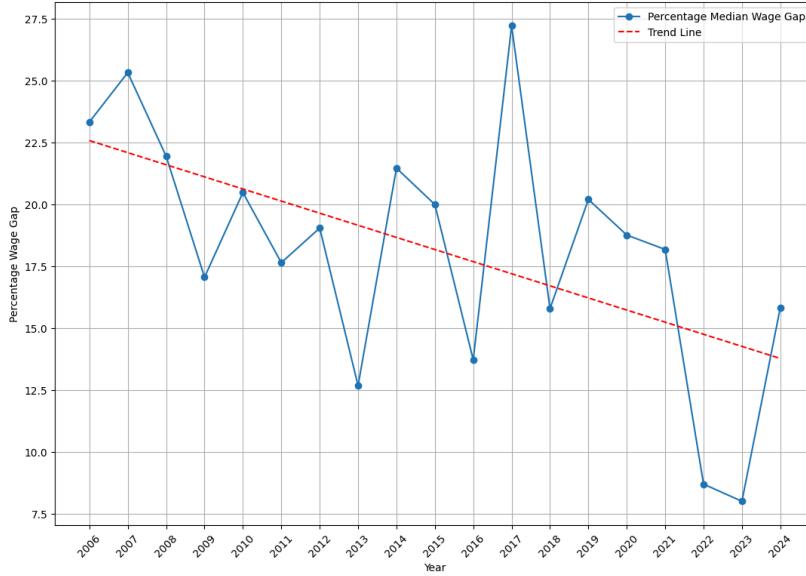
# plot percentage median wage gap by year
plt.figure(figsize=(13, 9))
plt.plot(percentage_median_wage_gap_by_year.index, percentage_median_wage_gap_by_year.values, marker='o', linestyle='--', label='Percentage Median Wage Gap')

# Add a trend line
z = np.polyfit(percentage_median_wage_gap_by_year.index, percentage_median_wage_gap_by_year.values, 1)
p = np.poly1d(z)
plt.plot(percentage_median_wage_gap_by_year.index, p(percentage_median_wage_gap_by_year.index), "r--", label='Trend Line')

plt.xlabel('Year')
plt.ylabel('Percentage Wage Gap')
plt.title('Percentage Median Wage Gap between Men and Women in Ontario by Year')
plt.xticks(percentage_median_wage_gap_by_year.index, rotation=45) # Set ticks for every year
plt.grid(True) # Add a grid for better readability
plt.legend()
plt.show()

```

Percentage Median Wage Gap between Men and Women in Ontario by Year



With the calculated percentage differences of the mean and median wage of Men and Women, we can see the existing gap appears to be declining over time.

```
# Apply a linear regression to percentage_mean_wage_gap_by_year
import statsmodels.api as sm

# Prepare data for regression
# The index of the series is the YEAR, which will be our independent variable
X = percentage_mean_wage_gap_by_year.index.values.reshape(-1, 1)
y = percentage_mean_wage_gap_by_year.values

# Add a constant to the independent variable for the intercept
X = sm.add_constant(X)

# Fit the linear regression model
model_gap_mean = sm.OLS(y, X).fit()

# Print the regression results
print(model_gap_mean.summary())

OLS Regression Results
=====
Dep. Variable: y R-squared: 0.046
Model: OLS Adj. R-squared: -0.010
Method: Least Squares F-statistic: 0.8274
Date: Mon, 10 Nov 2025 Prob (F-statistic): 0.376
Time: 19:53:58 Log-likelihood: -34.493
No. Observations: 19 AIC: 72.99
DF Residuals: 17 BIC: 74.88
DF Model: 1
Covariance Type: nonrobust
=====
coef std err t P>|t| [0.025 0.975]
-----
const 137.6880 132.643 1.038 0.314 -142.164 417.540
x1 -0.6599 0.866 -0.910 0.376 -0.199 0.079
-----
Omnibus: 7.597 Durbin-Watson: 2.136
Prob(Omnibus): 0.023 Jarque-Bera (JB): 5.023
Skew: 0.858 Prob(JB): 0.0812
Kurtosis: 4.859 Cond. No. 7.41e+05
-----
Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 7.41e+05. This might indicate that there are
strong multicollinearity or other numerical problems.
```

#### Observations

- R-squared is 0.046; approximately 4.6% of the variation can be explained by "Gender" and "Year"
- f-statistic is 0.8274 and prob (f-statistic) / p-value is 0.376, suggesting the regression is not statistically significant
- Intercept (const) = 137.688, p = 0.314, since p > 0.05, the regression is not statistically significant

```
# Plot the trend line for percentage_mean_wage_gap_by_year
# Get the parameters from the fitted model (from cell 2c3ae57e)
intercept = model_gap_mean.params[0] # The constant (intercept) is the first parameter
slope = model_gap_mean.params[1] # The coefficient for YEAR is the second parameter

# Get the years from the index of the series
years = percentage_mean_wage_gap_by_year.index

# Calculate the predicted values (trend line)
predicted_gap = intercept + slope * years

# Plot the original data points
plt.figure(figsize=(10, 6))
plt.plot(years, percentage_mean_wage_gap_by_year.values, marker='o', linestyle='-', label='Percentage Mean Wage Gap')

# Plot the trend line
plt.plot(years, predicted_gap, color='red', linestyle='--', label='Trend Line')

# Add labels and title
plt.xlabel('Year')
plt.ylabel('Percentage Wage Gap')
plt.title('Trend of Percentage Mean Wage Gap between Men and Women in Ontario')
plt.xticks(years, rotation=45) # Set ticks for every year
plt.grid(True) # Add a grid for better readability
plt.legend()
plt.show()
```

Trend of Percentage Mean Wage Gap between Men and Women in Ontario



```
# apply linear regression to percentage_median_wage_gap_by_year
# Prepare data for regression
# The index of the series is the YEAR, which will be our independent variable
X_median = percentage_median_wage_gap_by_year.index.values.reshape(-1, 1)
y_median = percentage_median_wage_gap_by_year.values

# Add a constant to the independent variable for the intercept
X_median = sm.add_constant(X_median)

# Fit the linear regression model
model_gap_median = sm.OLS(y_median, X_median).fit()

# Print the regression results
print(model_gap_median.summary())

OLS Regression Results
=====
Dep. Variable: y R-squared: 0.300
```

```

Model: OLS Adj. R-squared: 0.259
Method: Least Squares F-statistic: 7.294
Date: Mon, 10 Nov 2025 Prob (F-statistic): 0.0152
Time: 19:53:57 Log-Likelihood: -53.736
No. Observations: 19 AIC: 111.5
Df Residuals: 17 BIC: 113.4
Df Model: 1
Covariance Type: nonrobust
=====
```

Notes:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 7.41e+05. This might indicate that there are strong multicollinearity or other numerical problems.

#### Observations

- R-squared is 0.300; the model explains approximately 30% of the variation
- f-statistic is 7.294 and prob (f-statistic) / p-value is 0.0152, suggesting the model is statistically significant
- Intercept (const) = 1004.49, p = 0.014, since p < 0.05, the regression is statistically significant

```

# plot percentage median wage gap by year
# Get the parameters from the fitted model (from cell e85xiResDrU)
intercept = model_gap_median.params[0] # The constant (Intercept) is the first parameter
slope = model_gap_median.params[1] # The coefficient for YEAR is the second parameter

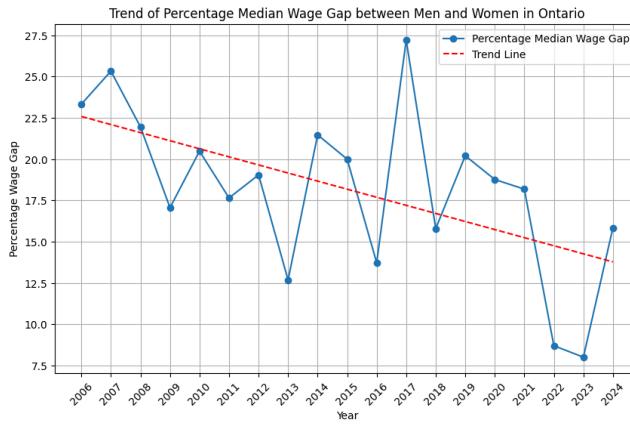
# Get the years from the index of the series
years = percentage_median_wage_gap_by_year.index

# Calculate the predicted values (trend line)
predicted_gap = intercept + slope * years

# Plot the original data points
plt.figure(figsize=(10, 6))
plt.plot(years, percentage_median_wage_gap_by_year.values, marker='o', linestyle='-', label='Percentage Median Wage Gap')

# Plot the trend line
plt.plot(years, predicted_gap, color='red', linestyle='--', label='Trend Line')

# Add labels and title
plt.xlabel('Year')
plt.ylabel('Percentage Wage Gap')
plt.title('Trend of Percentage Median Wage Gap between Men and Women in Ontario')
plt.xticks(years, rotation=45) # Set ticks for every year
plt.grid(True) # Add a grid for better readability
plt.legend()
plt.show()
```



From the linear regression model, both the mean and median percentage wage gap appear to be declining, meaning the gap seems to be narrowing.

The median percentage wage gap seems to be narrowing at a greater rate than the mean percentage.

Validate our model using K-Fold Cross-Validation.

```

# Use K-fold cross-validation to validate regression model of mean wage gap
from sklearn.model_selection import KFold
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import TimeSeriesSplit # Import TimeSeriesSplit

# Prepare data for regression (from cell 2cae57e)
X = percentage_mean_wage_gap_by_year.index.values.reshape(-1, 1)
y = percentage_mean_wage_gap_by_year.values

# Initialize the linear regression model
model = LinearRegression()

# Define the number of splits for time series cross-validation
n_splits = 5 # You can change the number of splits as needed
tscv = TimeSeriesSplit(n_splits=n_splits) # Use TimeSeriesSplit

# Perform Time Series cross-validation
# We will use the negative mean squared error as the scoring metric
# cross_val_score returns negative values for error metrics, so we take the absolute value
scores = cross_val_score(model, X, y, scoring='neg_mean_squared_error', cv=tscv)
mse_scores = -scores

# Print the results
print("Cross-validation results (Mean Squared Error) with {} splits:".format(n_splits))
print("Individual MSE scores: {}".format(mse_scores))
print("Mean MSE: {:.4f}".format(np.mean(mse_scores)))
print("Standard deviation of MSE: {:.4f}".format(np.std(mse_scores)))

Cross-validation results (Mean Squared Error) with 5 splits:
Individual MSE scores: [7.22586121 4.35227065 6.77123388 2.09175509 2.85441734]
Mean MSE: 4.6591
Standard deviation of MSE: 2.0490
```

#### Observations

- Mean MS is 4.6591
- Standard deviation of MSE: 2.0490

```

# Use K-fold cross-validation to validate regression model of median wage gap
# Prepare data for regression (from cell e85xiResDrU)
X_median = percentage_median_wage_gap_by_year.index.values.reshape(-1, 1)
y_median = percentage_median_wage_gap_by_year.values

# Initialize the linear regression model
model_median = LinearRegression()

# Define the number of splits for time series cross-validation
n_splits = 5 # You can change the number of splits as needed
tscv_median = TimeSeriesSplit(n_splits=n_splits) # Use TimeSeriesSplit

# Perform Time Series cross-validation
# We will use the negative mean squared error as the scoring metric
# cross_val_score returns negative values for error metrics, so we take the absolute value
scores_median = cross_val_score(model_median, X_median, y_median, scoring='neg_mean_squared_error', cv=tscv_median)
mse_scores_median = -scores_median

# Print the results
print("Cross-validation results (Mean Squared Error) for median wage gap with {} splits:".format(n_splits))
print("Individual MSE scores: {}".format(mse_scores_median))
print("Mean MSE: {:.4f}".format(np.mean(mse_scores_median)))
print("Standard deviation of MSE: {:.4f}".format(np.std(mse_scores_median)))

Cross-validation results (Mean Squared Error) for median wage gap with 5 splits:
Individual MSE scores: [26.62789623 27.79295166 45.10809607 4.93346516 56.66992236]
Mean MSE: 32.2265
Standard deviation of MSE: 17.6597
```

#### Observations

- Mean MS is 32.2265
- Standard deviation of MSE: 17.6597

Both models have moderate accuracy, however, they are both not consistently reliable. As the data set we are working with is aggregated (calculated mean and median difference), resulting in a smaller sample size, this may be causing the models to be unreliable.

### Research Question 3

Examine wage trends across the different subgroups between genders to determine influence of immigration status on wage gap  
Method: Apply decision tree model to analyze gender wage gaps across the different immigrant subgroups and validate model using K-Fold Cross-Validation

```

from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# Filter out suppressed wage values for both men and women
df_filtered = df_ontario[
    (df_ontario['Men_Wage_is_suppressed'] == False) &
    (df_ontario['Women_Wage_is_suppressed'] == False)
].copy()

# Create a 'wage_gap' target variable
df_filtered['wage_gap'] = df_filtered['Men_Wage'] - df_filtered['Women_Wage']

# Select features
features = ['YEAR', 'IMMIGRANT', 'Men_Wage', 'Women_Wage']
target = 'wage_gap'

X = df_filtered[features]
y = df_filtered[target]

# Identify categorical columns for encoding
categorical_features = ['IMMIGRANT'] # Only 'IMMIGRANT' is categorical in the updated features list

# Create a column transformer for one-hot encoding
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ],
    remainder='passthrough' # Keep other columns ('YEAR', 'Men_Wage', 'Women_Wage') as they are
)

# Create a pipeline with the preprocessor and the Decision Tree Regressor
pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                           ('regressor', DecisionTreeRegressor(random_state=42))])

# Define the number of splits for K-Fold Cross-Validation
n_splits = 5 # You can change the number of splits as needed
kf = KFold(n_splits=n_splits, shuffle=True, random_state=42) # Use KFold for general cross-validation

# Perform K-Fold Cross-Validation
# We will use the negative mean squared error as the scoring metric
scores = cross_val_score(pipeline, X, y, scoring='neg_mean_squared_error', cv=kf)
mse_scores = -scores

# Print the results
print("Cross-validation results (Mean Squared Error) for Decision Tree Regressor with {} folds:".format(n_splits))
print("Individual MSE scores: {}".format(mse_scores))
print("Mean MSE: {:.4f}".format(np.mean(mse_scores)))
print("Standard deviation of MSE: {:.4f}".format(np.std(mse_scores)))

# You can also fit the model on the entire dataset to examine feature importances
pipeline.fit(X, y)

# Get feature importances from the trained model
# Need to get the feature names after one-hot encoding
encoded_feature_names = pipeline.named_steps['preprocessor'].named_transformers_['cat'].get_feature_names_out(categorical_features)
passthrough_features = [col for col in features if col not in categorical_features]
all_feature_names = np.concatenate([encoded_feature_names, passthrough_features])

feature_importances = pd.Series(pipeline.named_steps['regressor'].feature_importances_, index=all_feature_names)
sorted_feature_importances = feature_importances.sort_values(ascending=False)

print("\nFeature Importances (Top 10) for predicting Wage Gap:")
print(sorted_feature_importances.head(10))

Cross-validation results (Mean Squared Error) for Decision Tree Regressor with 5 folds:
Individual MSE scores: [0.07235922 0.05322654 0.03345896 0.08072839 0.06082875]
Mean MSE: 0.0601
Standard deviation of MSE: 0.0163

Feature Importances (Top 10) for predicting Wage Gap:
Men_Wage          0.547634
Women_Wage         0.450986
YEAR              0.008995
IMMIGRANT_Total Landed Immigrants 0.00201
IMMIGRANT_Non-landed immigrants 0.000099
IMMIGRANT_Established immigrants, 10+ years 0.000056
IMMIGRANT_Recent immigrants, 5+ to 10 years 0.000044
IMMIGRANT_Born in Canada 0.000032
IMMIGRANT_Recent immigrants 5+ years 0.000025
IMMIGRANT_Total 0.000025
dtype: float64

```

The decision tree model shows very small numbers for all the immigrant categories, suggesting negligible impact on wages. The value is not 0, but the immigrant categories don't seem to add predictive value. The K-fold cross-validation seems suggests the model is moderately accurate and reliable.

### Research Question 4

Verify the correlation between education levels and wages between genders.

Method: Conduct a correlation analysis using Pearson correlation between education level and wages for each gender.

Use ANCOVA (analysis of covariance) to verify the effects of education levels on wages between genders.

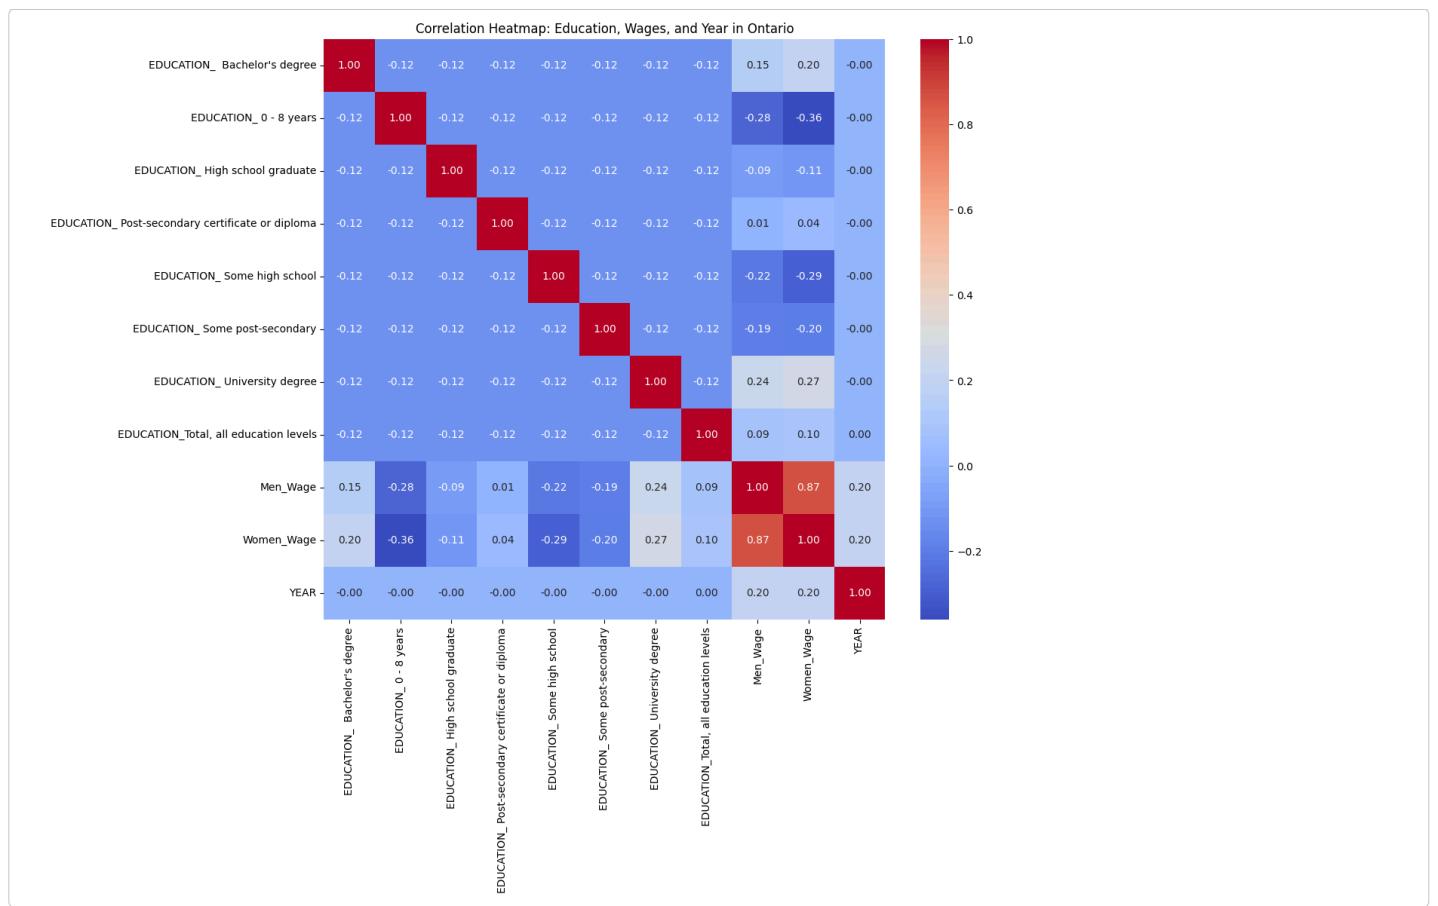
	EDUCATION_Bachelor's degree	EDUCATION_0 - 8 years	EDUCATION_High school graduate	EDUCATION_Post-secondary certificate or diploma	EDUCATION_Some high school	EDUCATION_Some post-secondary	EDUCATION_University degree	EDUCATION_Total, all education levels	Men_Wage	Women_Wage	YEAR
EDUCATION_Bachelor's degree	1.000000e+00	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	0.146492	0.196122	-2.102224e-16
EDUCATION_0 - 8 years	-1.250000e-01	1.000000e+00	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-0.277711	-0.360299	-1.060560e-15
EDUCATION_High school graduate	-1.250000e-01	-1.250000e-01	1.000000e+00	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-0.090160	-0.109367	-6.382926e-16
EDUCATION_Post-secondary certificate or diploma	-1.250000e-01	-1.250000e-01	-1.250000e-01	1.000000e+00	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	0.006639	0.036347	-8.347039e-16
EDUCATION_Some high school	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	1.000000e+00	-1.250000e-01	-1.250000e-01	-1.250000e-01	-0.219938	-0.287164	-1.118269e-15
EDUCATION_Some post-secondary	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	1.000000e+00	-1.250000e-01	-1.250000e-01	-0.191635	-0.202056	-1.931361e-16
EDUCATION_University degree	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	1.000000e+00	-1.250000e-01	0.235688	0.271631	-6.369096e-16
EDUCATION_Total, all education levels	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	-1.250000e-01	1.000000e+00	0.088542	0.096583	3.195423e-16
Men_Wage	1.464917e-01	-2.777110e-01	-9.015991e-02	6.639047e-03	-2.199384e-01	-1.916353e-01	2.356885e-01	8.854166e-02	1.000000	0.871257	2.020610e-01
Women_Wage	1.961218e-01	-3.602992e-01	-1.093668e-01	3.634654e-02	-2.871640e-01	-2.020557e-01	2.716306e-01	9.658311e-02	0.871257	1.000000	2.006960e-01
YEAR	-2.102224e-16	-1.060560e-15	-6.382926e-16	-8.347039e-16	-1.182690e-15	-1.931361e-16	-6.369096e-16	3.195423e-16	0.202061	0.200696	1.000000e+00

Next steps: [Generate code with correlation\\_matrix\\_education\\_wages\\_year](#) [New interactive sheet](#)

```

# Generate a heatmap of the education correlation matrix
plt.figure(figsize=(12, 10)) # Adjust the figure size for better readability
sns.heatmap(correlation_matrix_education_wages_year, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap: Education, Wages, and Year in Ontario')
plt.show()

```



Based on the correlation\_matrix\_education\_wages\_year:

- There is a positive correlation between having a 'University degree' and both 'Men\_Wage' (0.24) and 'Women\_Wage' (0.27)
- 'Bachelor's degree' shows a positive correlation with both 'Men\_Wage' (0.15) and 'Women\_Wage' (0.20), although slightly lower

This suggests that individuals with a university degree or Bachelor's degree tend to have higher wages.

- There is a negative correlation between having education '0 - 8 years' with both 'Men\_Wage' (-0.28) and 'Women\_Wage' (-0.36)
- There is a negative correlation also between having 'Some high school' with both 'Men\_Wage' (-0.22) and 'Women\_Wage' (-0.29)

This suggests that individuals with lower levels of education tend to have lower wages.

Other education levels like 'High school graduate', 'Post-secondary certificate or diploma', and 'Some post-secondary' show weaker correlations with wages, both positive and negative, depending on the specific category and gender.

The strength of the correlation seems to differ slightly. For example, having a 'University degree' appears to have a slightly stronger positive correlation with 'Women\_Wage' (0.27) compared to 'Men\_Wage' (0.24). Conversely, '0 - 8 years' of education has a stronger negative correlation with 'Women\_Wage' (-0.36) compared to 'Men\_Wage' (-0.28).

```

import statsmodels.api as sm
from statsmodels.formula.api import ols

# Prepare data for ANCOVA
# We will perform separate ANCOVA for Men's Wage and Women's Wage, with Education as the covariate

# Ensure 'EDUCATION' is treated as a categorical variable for the model formula
df_ontario['EDUCATION'] = df_ontario['EDUCATION'].astype('category')

# ANCOVA for Men's Wage with Education as covariate
# Formula: Men_Wage ~ C(EDUCATION)
# C() is used to treat EDUCATION as a categorical variable
model_men_ancova = ols('Men_Wage ~ C(EDUCATION)', data=df_ontario).fit()
anova_table_men = sm.stats.anova_lm(model_men_ancova, typ=2) # Use type 2 ANOVA for unbalanced designs

print("\nANCOVA Results for Men's Wage with Education as Covariate:")
print(anova_table_men)

# ANCOVA for Women's Wage with Education as covariate
# Formula: Women_Wage ~ C(EDUCATION)
model_women_ancova = ols('Women_Wage ~ C(EDUCATION)', data=df_ontario).fit()
anova_table_women = sm.stats.anova_lm(model_women_ancova, typ=2) # Use type 2 ANOVA for unbalanced designs

print("\nANCOVA Results for Women's Wage with Education as Covariate:")
print(anova_table_women)

# ANCOVA Results for Men's Wage with Education as Covariate:
# sum_sq   df      F PR(>F)
C(EDUCATION) 9.457823e+05  8.0  2174.466595  0.0
Residual     2.124791e+00  20511.0    NaN    NaN

# ANCOVA Results for Women's Wage with Education as Covariate:
# sum_sq   df      F PR(>F)
C(EDUCATION) 9.768410e+05  8.0  2174.466595  0.0
Residual     1.151776e+00  20511.0    NaN    NaN
/tmppython-input-1310997702.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy.
df_ontario['EDUCATION'] = df_ontario['EDUCATION'].astype('category')

```